

**Please execute the code below and observe the output you get. Also, please learn how to use each of these statements to get a similar task done.**

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
df = pd.read_csv('https://raw.githubusercontent.com/jackiekazil/data-wrangling/master/data/
(https://raw.githubusercontent.com/jackiekazil/data-wrangling/master/data) a/chp3/data-text.csv') df.head(2) df1 =
pd.read_csv('https://raw.githubusercontent.com/kjam/data-wranglingpycon/master/d
(https://raw.githubusercontent.com/kjam/data-wranglingpycon/master/d) ata/berlin_weather_oldest.csv')
df1.head(2)
```

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

In [43]:

```
df = pd.read_csv('https://raw.githubusercontent.com/jackiekazil/data-wrangling/master/data/
```

## 1. Get the Metadata from the above files.

In [44]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4656 entries, 0 to 4655
Data columns (total 12 columns):
Indicator                4656 non-null object
PUBLISH STATES           4656 non-null object
Year                    4656 non-null int64
WHO region               4656 non-null object
World Bank income group  4656 non-null object
Country                 4656 non-null object
Sex                     4656 non-null object
Display Value            4656 non-null int64
Numeric                 4656 non-null float64
Low                     0 non-null float64
High                    0 non-null float64
Comments                 0 non-null float64
dtypes: float64(4), int64(2), object(6)
memory usage: 436.6+ KB
```

## 2. Get the row names from the above files.

In [45]:

```
df.index.values
```

Out[45]:

```
array([ 0, 1, 2, ..., 4653, 4654, 4655], dtype=int64)
```

### 3. Change the column name from any of the above file.

In [46]:

```
df.rename(columns = {'Year':'years'})
```

...

In [47]:

```
df.head(1)
```

Out[47]:

	Indicator	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN

### 4. Change the column name from any of the above file and store the changes made permanently.

In [48]:

```
df.rename(columns = {'Indicator':'Indicator_id'},inplace=True)
```

In [49]:

```
df.head(1)
```

Out[49]:

	Indicator_id	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN

## 5. Change the names of multiple columns.

In [56]:

```
df.rename(columns = {'Country':'country','Sex':'sex'})
```

...

## 6. Arrange values of a particular column in ascending order.

In [58]:

```
df.sort_values(by = 'Year')
```

Out[58]:

	Indicator_id	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN
1270	Life expectancy at birth (years)	Published	1990	Europe	High-income	Germany	Male	72	72.0	NaN

## 7. Arrange multiple column values in ascending order.

In [59]:

```
df.sort_values(by = ['Year', 'Country'])
```

	at age 60 (years)			Mediterranean	income					
					group					
299	Life expectancy at birth (years)	Published	1990	Europe	Lower-middle-income	Albania	Male	67	67.0	Na
689	Life expectancy at birth (years)	Published	1990	Europe	Lower-middle-income	Albania	Both sexes	69	69.0	Na
1522	Life expectancy at age 60 (years)	Published	1990	Europe	Lower-middle-income	Albania	Both sexes	16	16.0	Na

## 8. Make country as the first column of the dataframe.

In [64]:

```
df.set_index('Country').reset_index()
```

		at age 60 (years)				income	sexes			
4	United Arab Emirates	Life expectancy at birth (years)	Published	2012	Eastern Mediterranean	High-income	Female	78	78.0	Na
5	Antigua and Barbuda	Life expectancy at birth (years)	Published	2000	Americas	High-income	Male	72	72.0	Na

## 10. Get the subset rows 11, 24, 37

In [86]:

```
df.iloc[[11,24,37]]
```

Out[86]:

	Indicator_id	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeri
11	Life expectancy at birth (years)	Published	2012	Europe	High-income	Austria	Female	83	83.0
24	Life expectancy at age 60 (years)	Published	2012	Western Pacific	High-income	Brunei Darussalam	Female	21	21.0
37	Life expectancy at age 60 (years)	Published	2012	Europe	High-income	Cyprus	Female	26	26.0

## 11. Get the subset rows excluding 5, 12, 23, and 56

In [95]:

```
d1=df.index.isin([5,12,23,56])
df[~d1]
```

	at age 60 (years)				income					
3	Life expectancy at age 60 (years)	Published	2000	Europe	High-income	Andorra	Both sexes	23	23.0	Na
4	Life expectancy at birth (years)	Published	2012	Eastern Mediterranean	High-income	United Arab Emirates	Female	78	78.0	Na
6	Life expectancy at age 60 (years)	Published	1990	Americas	High-income	Antigua and Barbuda	Male	17	17.0	Na
7	Life expectancy at birth (years)	Published	2012	Europe	High-income	Austria	Female	83	83.0	Na

## Load datasets from CSV

```
users= pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/users.csv')
(https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/users.csv')
```

```
sessions = pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/sessions.csv')
(https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/sessions.csv')
```

```
products = pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/products.csv')
(https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/products.csv')
```

```
transactions =
pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/transactions.csv')
(https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/transactions.csv')
```

In [3]:

```
users= pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/user
sessions = pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/
products = pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/
transactions = pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/D
```

## 12. Join users to transactions, keeping all rows from transactions and only matching rows from users (left join)

In [17]:

```
transactions.join(users,on = 'UserID',how='left',lsuffix='_transactions',rsuffix='_users')
```

Out[17]:

	TransactionID	TransactionDate	UserID_transactions	ProductID	Quantity	UserID_us
0	1	2010-08-21	7	2	1	NaN
1	2	2011-05-26	3	4	1	4
2	3	2011-06-16	3	3	1	4
3	4	2012-08-26	1	2	3	2
4	5	2013-06-06	2	4	1	3
5	6	2013-12-23	2	5	6	3
6	7	2013-12-30	3	4	1	4
7	8	2014-04-24	NaN	2	3	NaN
8	9	2015-04-24	7	4	3	NaN
9	10	2016-05-08	3	4	4	4

In [20]:

```
transactions.merge(users,'left',on='UserID')
```

Out[20]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	User	Gender	Regis
0	1	2010-08-21	7	2	1	NaN	NaN	NaN
1	2	2011-05-26	3	4	1	Caroline	female	2012-
2	3	2011-06-16	3	3	1	Caroline	female	2012-
3	4	2012-08-26	1	2	3	Charles	male	2012-
4	5	2013-06-06	2	4	1	Pedro	male	2010-
5	6	2013-12-23	2	5	6	Pedro	male	2010-
6	7	2013-12-30	3	4	1	Caroline	female	2012-
7	8	2014-04-24	NaN	2	3	NaN	NaN	NaN
8	9	2015-04-24	7	4	3	NaN	NaN	NaN
9	10	2016-05-08	3	4	4	Caroline	female	2012-

In [45]:

```
import sqlite3 as db
```

In [46]:

```
conn = db.connect('test.db')
```

In [121]:

```
transactions.to_sql('trxn',conn,if_exists="replace")
users.to_sql('users',conn,if_exists="replace")
products.to_sql('products',conn,if_exists="replace")
sessions.to_sql('sessions',conn,if_exists="replace")
```

### 13. Which transactions have a UserID not in users?

In [53]:

```
pd.read_sql_query('select * from txn where userID not in (select userID from users) or use
```

Out[53]:

	index	TransactionID	TransactionDate	UserID	ProductID	Quantity
0	0	1	2010-08-21	7.0	2	1
1	7	8	2014-04-24	NaN	2	3
2	8	9	2015-04-24	7.0	4	3

In [131]:

```
transactions[~transactions['UserID'].isin(users['UserID'])]
```

Out[131]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	ranked
0	1	2010-08-21	7.0	2	1	1.0
7	8	2014-04-24	NaN	2	3	NaN
8	9	2015-04-24	7.0	4	3	2.0

## 14. Join users to transactions, keeping only rows from transactions and users that match via UserID (inner join)

In [56]:

```
pd.read_sql_query('''
select t.* from txn as t inner join users as u
on u.userID=t.userID''', conn)
```

Out[56]:

	index	TransactionID	TransactionDate	UserID	ProductID	Quantity
0	1	2	2011-05-26	3.0	4	1
1	2	3	2011-06-16	3.0	3	1
2	3	4	2012-08-26	1.0	2	3
3	4	5	2013-06-06	2.0	4	1
4	5	6	2013-12-23	2.0	5	6
5	6	7	2013-12-30	3.0	4	1
6	9	10	2016-05-08	3.0	4	4



In [132]:

```
transactions.merge(users, how='inner', on='UserID')
```

Out[132]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	ranked	User	Gender
0	2	2011-05-26	3	4	1	1.0	Caroline	female
1	3	2011-06-16	3	3	1	2.0	Caroline	female
2	7	2013-12-30	3	4	1	3.0	Caroline	female
3	10	2016-05-08	3	4	4	4.0	Caroline	female
4	4	2012-08-26	1	2	3	1.0	Charles	male
5	5	2013-06-06	2	4	1	1.0	Pedro	male
6	6	2013-12-23	2	5	6	2.0	Pedro	male

**15. Join users to transactions, displaying all matching rows AND all non-matching rows (full outer join)**

In [60]:

```
transactions.merge(users, 'outer', on='UserID')
```

Out[60]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	User	Gender	Reg
0	1.0	2010-08-21	7.0	2.0	1.0	NaN	NaN	NaN
1	9.0	2015-04-24	7.0	4.0	3.0	NaN	NaN	NaN
2	2.0	2011-05-26	3.0	4.0	1.0	Caroline	female	201
3	3.0	2011-06-16	3.0	3.0	1.0	Caroline	female	201
4	7.0	2013-12-30	3.0	4.0	1.0	Caroline	female	201
5	10.0	2016-05-08	3.0	4.0	4.0	Caroline	female	201
6	4.0	2012-08-26	1.0	2.0	3.0	Charles	male	201
7	5.0	2013-06-06	2.0	4.0	1.0	Pedro	male	201
8	6.0	2013-12-23	2.0	5.0	6.0	Pedro	male	201
9	8.0	2014-04-24	NaN	2.0	3.0	NaN	NaN	NaN
10	NaN	NaN	4.0	NaN	NaN	Brielle	female	201
11	NaN	NaN	5.0	NaN	NaN	Benjamin	male	201

## 16. Determine which sessions occurred on the same day each user registered

In [65]:

```
pd.read_sql_query('''select * from users as u inner join sessions as s
on u.userID=s.userID
and SessionDate = Registered
''', conn)
```

Out[65]:

	index	UserID	User	Gender	Registered	Cancelled	index	SessionID	SessionDate
0	1	2	Pedro	male	2010-08-01	2010-08-08	1	2	2010-08-01
1	3	4	Brielle	female	2013-07-17	None	8	9	2013-07-17

In [133]:

```
pd.merge(left=users, right=sessions, how='inner', left_on=['UserID', 'Registered'], right_c
```

Out[133]:

	UserID	User	Gender	Registered	Cancelled	SessionID	SessionDate
0	2	Pedro	male	2010-08-01	2010-08-08	2	2010-08-01
1	4	Brielle	female	2013-07-17	NaN	9	2013-07-17

## 17. Build a dataset with every possible (UserID, ProductID) pair (cross join)

In [68]:

```
pd.read_sql_query('''select u.userID,p.ProductID from users as u join products as p  
''', conn)
```

Out[68]:

	<b>UserID</b>	<b>ProductID</b>
<b>0</b>	1	1
<b>1</b>	1	2
<b>2</b>	1	3
<b>3</b>	1	4
<b>4</b>	1	5
<b>5</b>	2	1
<b>6</b>	2	2
<b>7</b>	2	3
<b>8</b>	2	4
<b>9</b>	2	5
<b>10</b>	3	1
<b>11</b>	3	2
<b>12</b>	3	3
<b>13</b>	3	4
<b>14</b>	3	5
<b>15</b>	4	1
<b>16</b>	4	2
<b>17</b>	4	3
<b>18</b>	4	4
<b>19</b>	4	5
<b>20</b>	5	1
<b>21</b>	5	2
<b>22</b>	5	3
<b>23</b>	5	4
<b>24</b>	5	5

In [134]:

```
df1 = pd.DataFrame({'key': np.repeat(1, users.shape[0]), 'UserID': users.UserID})
df2 = pd.DataFrame({'key': np.repeat(1, products.shape[0]), 'ProductID': products.ProductID})
pd.merge(df1, df2, on='key')[['UserID', 'ProductID']]
```

Out[134]:

	<b>UserID</b>	<b>ProductID</b>
<b>0</b>	1	1
<b>1</b>	1	2
<b>2</b>	1	3
<b>3</b>	1	4
<b>4</b>	1	5
<b>5</b>	2	1
<b>6</b>	2	2
<b>7</b>	2	3
<b>8</b>	2	4
<b>9</b>	2	5
<b>10</b>	3	1
<b>11</b>	3	2
<b>12</b>	3	3
<b>13</b>	3	4
<b>14</b>	3	5
<b>15</b>	4	1
<b>16</b>	4	2
<b>17</b>	4	3
<b>18</b>	4	4
<b>19</b>	4	5
<b>20</b>	5	1
<b>21</b>	5	2
<b>22</b>	5	3
<b>23</b>	5	4
<b>24</b>	5	5

**18. Determine how much quantity of each product was purchased by each user**

In [73]:

```
pd.read_sql_query('''
select userID,ProductID,sum(Quantity)
from txn as t
group by userID,ProductID
having t.userID is not null
''', conn)
```

Out[73]:

	<b>UserID</b>	<b>ProductID</b>	<b>sum(Quantity)</b>
<b>0</b>	1.0	2	3
<b>1</b>	2.0	4	1
<b>2</b>	2.0	5	6
<b>3</b>	3.0	3	1
<b>4</b>	3.0	4	6
<b>5</b>	7.0	2	1
<b>6</b>	7.0	4	3

In [136]:

```
df1 = pd.DataFrame({'key': np.repeat(1, users.shape[0]), 'UserID': users.UserID})
df2 = pd.DataFrame({'key': np.repeat(1, products.shape[0]), 'ProductID': products.ProductID})
user_products = pd.merge(df1, df2, on='key')[['UserID', 'ProductID']]
pd.merge(user_products, transactions, how='left', on=['UserID', 'ProductID']).groupby(['UserID', 'ProductID']).sum().reset_index()
```

Out[136]:

	<b>UserID</b>	<b>ProductID</b>	<b>Quantity</b>
<b>0</b>	1	1	0.0
<b>1</b>	1	2	3.0
<b>2</b>	1	3	0.0
<b>3</b>	1	4	0.0
<b>4</b>	1	5	0.0
<b>5</b>	2	1	0.0
<b>6</b>	2	2	0.0
<b>7</b>	2	3	0.0
<b>8</b>	2	4	1.0
<b>9</b>	2	5	6.0
<b>10</b>	3	1	0.0
<b>11</b>	3	2	0.0
<b>12</b>	3	3	1.0
<b>13</b>	3	4	6.0
<b>14</b>	3	5	0.0
<b>15</b>	4	1	0.0
<b>16</b>	4	2	0.0
<b>17</b>	4	3	0.0
<b>18</b>	4	4	0.0
<b>19</b>	4	5	0.0
<b>20</b>	5	1	0.0
<b>21</b>	5	2	0.0
<b>22</b>	5	3	0.0
<b>23</b>	5	4	0.0
<b>24</b>	5	5	0.0

## 19. For each user, get each possible pair of pair transactions (TransactionID1, TransactionID2)

Question is not clear

In [130]:

```
pd.merge(transactions, transactions, on='UserID')
```

Out[130]:

	TransactionID_x	TransactionDate_x	UserID	ProductID_x	Quantity_x	ranked_x	Tr
0	1	2010-08-21	7.0	2	1	1.0	1
1	1	2010-08-21	7.0	2	1	1.0	9
2	9	2015-04-24	7.0	4	3	2.0	1
3	9	2015-04-24	7.0	4	3	2.0	9
4	2	2011-05-26	3.0	4	1	1.0	2
5	2	2011-05-26	3.0	4	1	1.0	3
6	2	2011-05-26	3.0	4	1	1.0	7
7	2	2011-05-26	3.0	4	1	1.0	10
8	3	2011-06-16	3.0	3	1	2.0	2
9	3	2011-06-16	3.0	3	1	2.0	3
10	3	2011-06-16	3.0	3	1	2.0	7
11	3	2011-06-16	3.0	3	1	2.0	10
12	7	2013-12-30	3.0	4	1	3.0	2
13	7	2013-12-30	3.0	4	1	3.0	3
14	7	2013-12-30	3.0	4	1	3.0	7
15	7	2013-12-30	3.0	4	1	3.0	10
16	10	2016-05-08	3.0	4	4	4.0	2
17	10	2016-05-08	3.0	4	4	4.0	3
18	10	2016-05-08	3.0	4	4	4.0	7
19	10	2016-05-08	3.0	4	4	4.0	10
20	4	2012-08-26	1.0	2	3	1.0	4
21	5	2013-06-06	2.0	4	1	1.0	5
22	5	2013-06-06	2.0	4	1	1.0	6
23	6	2013-12-23	2.0	5	6	2.0	5
24	6	2013-12-23	2.0	5	6	2.0	6
25	8	2014-04-24	NaN	2	3	NaN	8

**20. Join each user to his/her first occuring transaction in the transactions table**



In [109]:

```
transactions['ranked']=transactions.groupby('UserID')['TransactionDate'].rank(ascending=1,n
```

In [110]:

```
df=transactions[transactions['ranked'] == 1]
```

In [114]:

```
users.merge(df,'left','UserID')
```

Out[114]:

	UserID	User	Gender	Registered	Cancelled	TransactionID	TransactionDate	Pr
0	1	Charles	male	2012-12-21	NaN	4.0	2012-08-26	2.0
1	2	Pedro	male	2010-08-01	2010-08-08	5.0	2013-06-06	4.0
2	3	Caroline	female	2012-10-23	2016-06-07	2.0	2011-05-26	4.0
3	4	Brielle	female	2013-07-17	NaN	NaN	NaN	NaN
4	5	Benjamin	male	2010-11-25	NaN	NaN	NaN	NaN

## 21. Test to see if we can drop columns

In [119]:

```
users.drop('Gender',axis=1)
```

Out[119]:

	UserID	User	Registered	Cancelled
0	1	Charles	2012-12-21	NaN
1	2	Pedro	2010-08-01	2010-08-08
2	3	Caroline	2012-10-23	2016-06-07
3	4	Brielle	2013-07-17	NaN
4	5	Benjamin	2010-11-25	NaN

In [ ]:

```
users.drop('Gender',axis=1,inplace=True)
```